

# Machine Learning & Subsymbolic Computing

Chapters 10 (to end of 10.2),  
11 (to end of 11.2)

*(disclaimer: most images are not mine)*

1

## Machine Learning

- ◆ Obviously learning is important to any domain
- ◆ Machine learning is often associated with lower-level systems where learning is most of what the system does: e.g. neural networks
- ◆ But learning components are important in any agent designed to work in a changing environment
- ◆ Learning allows such agents to adapt as the domain changes around them

2

## Symbolic vs. Subsymbolic Learning

- ◆ Learning can be done at both the symbolic and subsymbolic levels
- ◆ I'm first going to talk about learning in general, then learning in subsymbolic systems, and finally learning in symbolic systems
- ◆ Because you haven't seen any subsymbolic systems yet, this will also introduce you to this topic

3

## Learning

- ◆ The number of variables in learning is HUGE
- ◆ We can talk about WHAT is being learned for example: concepts? goodness of concepts? experiences to repeat or not to repeat?
- ◆ As you'd imagine, there's specialized techniques on that
- ◆ There's also ways of teaching, and indeed, degrees to which an instructor is even involved...


4

## Learning Method

- ◆ We have a spectrum of learning mechanisms based on how involved an external instructor is in the process
- ◆ Highest to lowest involvement:
- ◆ Rote learning
  - Direct implantation of knowledge without inference on the part of the learner
  - i.e. I tell you the answer, you remember it and recall it at the right time later on

5

## Learning Method



- ◆ Learning from Instruction
  - transformation of external instruction into an internal representation and integration with prior knowledge and skills

6

## Learning Method

- ◆ Learning from Example
  - Extraction and refinement of knowledge from specific cases – e.g. an arch

Positive Examples




Negative Examples



7

## Learning Method


- ◆ Learning by Analogy
  - Transferring a solution from one situation to a new one with novel aspects



(this is a very weak analogy!)

8

## Learning Method



- ◆ Learning by Discovery
  - Gathering new knowledge and skills from observation and experience
  - easy to see that this is well beyond the other levels!

9

## Learning Feedback

- ◆ Feedback indicates performance level achieved so far. We can differentiate learning based on the type of feedback given to an agent as well:
  - Supervised Learning: Feedback specifies desired activity of learner and objective of learning is to match this
  - Reinforcement Learning: Feedback only specifies utility of actual activity of the learner and objective is to maximize this utility
  - Unsupervised Learning: No explicit feedback provided, objective is to find useful and desired activities based on trial and error and self-organization

10

## Types of Reasoning

- ◆ You're well-used to deductive reasoning by now! We have some axiom, and use logical reasoning to DEDUCE some new axiom
- ◆ This deduction is certain and solid provided we have everything purely FOL based
- ◆ This isn't the only type of reasoning we can perform though
- ◆ We can and do use non-deductive reasoning

11

## Types of Reasoning

### ◆ Abduction

b	conclusion
$a \rightarrow b$	rule
<hr/>	
$\therefore a$	premise

- ◆ A could have caused B
- ◆ Recall from implication that this is not necessarily so – but we CAN reason about the possibility it caused it
- ◆ This is a major part of diagnostic reasoning

12

## Types of Reasoning

### ◆ Induction

a      premise  
b      conclusion

---

$\therefore a \rightarrow b$  rule

- ◆ This is a major part of reasoning in learning: associating concepts!
- ◆ Again, this MIGHT be wrong (not sound!), but we do it all the time!

13

## Symbolic Vs. Subsymbolic Learning

- ◆ In the case of symbolic learning, we also have to choose a representation for what's to be learned
- ◆ This is not the case in subsymbolic learning
- ◆ Why?
- ◆ There IS no representation
- ◆ Recall subsymbolic reasoning: it exists below the symbol level and does not construct symbolic representations

14

## Symbolic vs. Subsymbolic

- ◆ Instead of building and manipulating symbols:  
bird(tweety)  
If bird(X) Then hasWings(X)  
 $\therefore$  hasWings(tweety)
- ◆ We recognize (match) patterns in data without using symbols or maintaining symbol structures

15

## Cognitive Processing

- ◆ This is very fast (we recognize faces in <100ms, and this is not an easy task!)
- ◆ But also unexplainable – we can't say *WHY* we recognize a face when we're doing it subsymbolically, it just happens
- ◆ The lack of explainability emphasizes the lack of representation
- ◆ The overall SYSTEM recognizes faces, but no single part or aspect is stored anywhere in particular
  - this is why subsymbolic approaches can be very robust, but also very hard to understand

16

## Connectionism

- ◆ Most subsymbolic implementations involve connectionist approaches
- ◆ A connectionist approach is one that is centered around the connections between simple elements
- ◆ Connectionist approaches don't HAVE to be subsymbolic
  - For example, if I connect behaviours together so the stimulus of one leads to another, that's a connectionist concept...

17

## Subsymbolic AI

- ◆ Connectionist approaches are most commonly associated with low level subsymbolic problems, however
- ◆ This paradigm has been used very successfully to develop intelligent systems that perform low-level cognitive tasks such as speech recognition
- ◆ Neural networks are the most well known form of connectionist approach (there are many others!)
- ◆ Most artificial neural networks operate in a manner that is significantly simpler than that of actual biological neurons, but are based on the same concepts

18

## Neural Networks

- ◆ Neural Networks are pattern matching systems. Some researchers believe that we do not perform explicit symbol manipulation when solving many of the tasks that are considered to require intelligence
- ◆ Instead, we perform a pattern matching process in which we react to stimuli at a subconscious level instead of explicitly processing the stimuli at a conscious level


19

## Neural Networks

- ◆ Neural networks are models of intelligent processing that consist of large numbers of simple processing units that collectively are able to perform very complex pattern matching tasks
- ◆ The processor can perform only one type of operation; the processors are quite slow
- ◆ Doesn't sound like much but: Connections link the processors together
- ◆ This massive parallelism lets them do very interesting things

20

## Human Neural Cells




- ◆ There are approximately  $10^{11}$  neurons in the brain
- ◆ There are approximately  $10^{15}$  connections in the brain
- ◆ The brain performs about  $10^4$  operations per second

21

## Neural Networks


- ◆ Artificial neural networks attempt to duplicate the processing of neural cells: they're a physical model of intelligence rather than a logical one
- ◆ A unit "turns on" (dark connections in image below) if it receives sufficient stimuli (from the outside or from other neurons) and in turn stimulates others or generates output



22

## Neural Networks

- ◆ Neural networks are trained by presenting stimulus-response patterns and adjusting weights (strengths between connections) until the desired response is generated for each stimulus
- ◆ In this example, the network below associates short and thin with tweety – it's small enough you can see how the connections work




- ◆ but in any non-trivial network we could never artificially weight (i.e. hand construct) the connections to our desired outputs this way!

23

## Neural Networks

- ◆ Instead, we design learning algorithms that cause the weights in a network to be adjusted based on training cases provided to the network



- ◆ We could take patterns and get the network to adjust itself to the network on the previous slide. A different set of patterns could be used to train the network to associate cute, short, and thin with Tweety
- ◆ Clearly not just the patterns, but the information we give the network when it is incorrect (supervised learning)

24


## Processing

- ◆ The inputs to each node are multiplied by the strength (weight) of the connection to the node and then summed to produce the net input (or stimulus) to the node

$$\text{Net} = \sum I_i * W_i$$

## Processing


- ◆ The net input is then sent through an activation or transfer function that generates the output value (or response) for the node



- ◆ A typical transfer function is the threshold function
  - at a certain input threshold an output is given, and none is given if that threshold is not reached
- ◆ Using a Sigmoid function for transfer is also popular – what effect would this have?

## Training

- ◆ A neural network is trained by presenting patterns to the network
- ◆ For units that do not generate the desired output, the weights coming into that unit are adjusted slightly by a learning algorithm




- ◆ An error term is computed for each output unit
- ◆ The error term is a function of the difference between *the value actually generated by the unit* and *the value that was supposed to have been generated (supervised learning)*

# Single-Layer Networks

- ◆ Networks with direct connections from the input layer to the output layer are referred to as single-layer networks because there is only one layer of processing units
- ◆ Single-layer networks are very limited in their ability to perform stimulus-response mappings
- ◆ For example, single-layer networks cannot learn an exclusive-or mapping
- ◆ This was why perceptrons were initially abandoned

## Multi-Layer Networks

- ◆ Adding one or more additional layers of hidden units between the input layer and the output layer permits the network to perform **MUCH** more complex stimulus-response mappings



29

## Backpropagation Networks

- ◆ Backpropagation neural networks are the most popular type of multi-layer network
- ◆ Until relatively recently, it was not known how to train backpropagation networks
- ◆ Finding error and using it to alter weights on output nodes is easy – just difference between desired and actual outputs
- ◆ But what about hidden layers? How can we find out how much error some hidden node is responsible for? No way to know for sure...

30

## Backpropagation

- ◆ The error is instead approximated, then *propagated back* through hidden layers
- ◆ I won't state the approximation formula used, or even its characteristics here – to even discuss those would require getting into low level learning functions beyond the level of 3190 – You will see this in 4360 next term
- ◆ Error ends up being propagated back to the input nodes, and each node can adjust connection weights as it goes back

31

## Example: NetTalk (p. 471-3)

- ◆ Terry Sejnowski of Johns Hopkins developed a system that can pronounce words of text
- ◆ The system consists of a backpropagation network with 203 input units (29 text characters, 7 characters at a time), 80 hidden units, and 26 output units
  - The system was developed over a year
- ◆ The DEC-talk system consists of hand-coded linguistic rules for speech pronunciation
  - developed over approximately 10 years
- ◆ DEC-talk outperforms NETtalk but DEC-talk required significantly more development time

32



## NetTalk

- ◆ "This exemplifies the utility of neural networks; they are easy to construct and can be used even when a problem is not fully understood. However, rule-based algorithms usually out-perform neural networks when enough understanding is available"

– Hertz, *Introduction to the Theory of Neural Networks*, p. 133

33

## Remember

- ◆ What we're doing is strictly supervised learning here
- ◆ The system is learning patterns that **we** are organizing for it and presenting in a rational way, and we're telling it whether a new item fits the pattern or not
- ◆ The system can't explain WHY it feels something matches or not
  - again, typical of subsymbolic approaches
- ◆ This means the system can't explain what it's learning either...how are we satisfied that it is doing what we want in the end?

34

## Tank Recognition System

- ◆ Researchers trained a neural network to differentiate scenes that contained one or more tanks from scenes that did not contain tanks
- ◆ The network was able to correctly recognize new scenes that the network had not been trained with
- ◆ The network was finally tested with additional scenes and was not able to recognize them correctly
- ◆ diagnosis?

35

## Tank Recognition System

- ◆ The network actually recognized that the pictures taken of the tanks were taken on sunny days while the pictures taken without tanks were taken on cloudy days
- ◆ The system WAS learning – just not what we expected
  - This is mainly a problem with US not recognizing the features and not differentiating what is intended to be learned (a faulty teacher)
  - also really a common-sense reasoning problem (think of a person doing this)

36

## Symbolic Machine Learning

- ◆ We can (and do!) use learning at a level above the subsymbolic as well
- ◆ One common method is through the use of *Version Space Search*. The name comes from the fact that we are searching through the space of different versions of a concept as we see positive and negative examples
- ◆ Version space search uses two basic operations: generalization and specialization
- ◆ Generalization makes more general statements:  
color(ball,red) can become color(X,red)
  - size(small)^color(red) can become  
size(small) ^ (color(red)vcolour(blue))

37

## Version Space Search

- ◆ Similarly, specialization takes those more general statements and applies them more specifically (e.g. the inverse of the previous two examples!)
- ◆ What version space search does is attempt to find symbolic expressions that accurately describe a set of examples
- ◆ For example, various examples of an arch
- ◆ We can give a symbolic description of several arches and ask the system to search for a common description

38

## Version Space Search

- ◆ That search, with generalization, will allow it to make a general statement true of all arches it knows about
- ◆ as we add more examples, that description may prove too general
  - i.e. cover things that aren't arches
- ◆ the system is then expected to use specialization to find coverings that deal with the positive cases but exclude the negative ones

39

## Version Space Search

- ◆ One such approach to search: Specific to General
- ◆ We maintain a set of candidate concept definitions, and ensure they are maximally specific generalizations
- ◆ This means that a concept covers all positive examples, no negative examples, and is more specific than any other concept that does
- ◆ We start out extremely specific (just this one case is an arch!), and gradually generalize to include all valid candidates

40

## Algorithm

- ◆ S=set of candidate hypotheses; N is the set of all negative instances seen so far
- ◆ initialize S to be the 1<sup>st</sup> positive training instance (i.e. S starts off extremely specific – one case!)
- ◆ For every training instance {  
if the instance is a Positive instance, P {  
for every element s in S: {if s doesn't match P,  
replace s with its most specific generalizations  
that do match P} (there may be several of these)  
Delete from S any hypothesis more general than  
some other hypothesis in S  
Delete from S all hypotheses that match a  
previously observed negative instance in N } //if

(continued...)

41

## Algorithm

- ◆ If the instance is a negative instance, N {  
Delete all members of S that match n  
add n to N to check future hypotheses for  
overgeneralization } } //if, for
- ◆ This requires both positive and negative examples –  
negative examples keep the algorithm from  
overgeneralizing
- ◆ How else could we have done this?
- ◆ we could also have gone the other way – started very  
general, then got more specific
- ◆ i.e. start by saying EVERYTHING is an arch!
- ◆ Then go through negative examples making our definition  
of an arch more and more specific!

42

## Candidate Elimination

- ◆ What we'd be keeping track of is a set of concepts  
that are maximally general and still cover all positive  
instances and no negative ones
- ◆ We'd need positive instances here to avoid getting  
too specific
- ◆ Taken together, these form a BIDIRECTIONAL  
search that we call Candidate Elimination
- ◆ We maintain TWO sets of candidate concepts: the  
maximally general set and the maximally specific set
- ◆ we specialize G and generalize S until we converge  
on the target!

43

## Symbolic Machine Learning

- ◆ There are many other types of machine  
learning at the symbolic level
- ◆ For example, we can learn decision trees  
based on cases
- ◆ e.g. a loan manager looking at specific  
questions that were asked, and from that  
whether people later defaulted on loans
- ◆ We can integrate cases one at a time and  
gradually learn a tree that lets us decide a  
risk assessment for different categories

44

## Symbolic Machine Learning

- ◆ this is analogous to our version space scenario
- ◆ we're attempting to induce a decision tree using cases at hand, and alter it based on new positive and negative cases
- ◆ we add new branches or merge branches into similar decisions (specialize, generalize)
- ◆ We can again do this top down, bottom up, or bidirectionally (similar issues)
- ◆ There are still different types of learning yet –

45

## Reinforcement Learning

- ◆ In our robotics situation, we could learn a set of reactions based on experiences in an environment
- ◆ The previous situations were all examples of SUPERVISED LEARNING – the most basic type, where a teacher organizes the cases and presents them in a coherent fashion
- ◆ In robotics, we want them be more independent from the actions of a teacher – we're going to do REINFORCEMENT LEARNING
- ◆ i.e. they just get good/bad (from a teacher or something in the world, such as a score)

46

## Q-Learning

- ◆ We're going to look at a simplified version of Q-learning, a form of reinforcement learning
- ◆ Think of the robot having a big reaction table
- ◆ rows: Actions I can perform; cols: situations I can recognize (it's a mapping, doesn't really have to be a table)
- ◆ The table cells can contain RANKINGS of how good each potential action is for all situations
- ◆ we call this table Q (our Q-table)
- ◆ We could fill this with random values to start with
- ◆ it would then be a (very poor) approximation of  $Q^*$ , the optimal table for the environment
- ◆ A learning algorithm would allow the agent to improve its table

47

## Q-learning

- ◆ When it does something good (think scoring a goal in soccer) we'll give it positive reinforcement
  - Yay....
- ◆ when it does something bad, we'll give it negative reinforcement
  - Booo....
- ◆ and the learning algorithm will adjust the table's "goodness" entry for that action down or up
- ◆ (What kind of algorithm is this?)
  - Iterative Improvement!
- ◆ What action do we pick on each turn?

48

## The Best?

- ◆ No – we'll have local maximum problems
- ◆ Suppose my highest ranked action initially (i.e. randomly) is 8, and the next highest is 5
- ◆ We do the action, goodness gets adjusted to 7
- ◆ Which do we do next time?
- ◆ The 7, if we're choosing by goodness
- ◆ but what's the 5?
  - a *GUESS*..could really be a 9 (even worse, what if all untested actions were 0?)
- ◆ We need to be able to find out whether that 5 is *really* a 5 or not...so what do we do?

49

## Exploration vs. Exploitation

- ◆ We choose the best action *sometimes*, and some other choice at other times
- ◆ most often a RANDOM choice
  - Randomness is extremely important in learning, this is a good example
  - Being able to try anything randomly also assumes that all actions are safe to try – may want to initially bias/restrict some
- ◆ This also considers only the immediate reaction in dealing with reinforcement
- ◆ What's the problem with reinforcing individual reactions?
- ◆ Not often accurate!

50

## Delayed Reinforcement

- ◆ You might get delayed reinforcement
  - I work today, I get paid @ the end of the month...
  - reinforcement is naturally delayed sometimes – how can we learn from that?
- ◆ In multi-agent settings the reinforcement is delayed in other ways
  - You kick a ball, somebody else scores a goal – how do you know what you did was good?
  - or you screw up, somebody else scores a goal – how do you know what you did was bad?
- ◆ These are instances of the Credit Assignment Problem (former: *temporal* credit assignment; latter: *inter-agent* credit assignment)

51

## Dealing with the CAP

- ◆ need to deal with the fact that your immediate action didn't necessarily cause what's going on
- ◆ we do this in Q-learning by passing a portion of the blame/praise *back to previous actions*
- ◆ Divide reinforcement up into two parts, I and P – I reinforces the immediate action, P reinforces the previous one
- ◆ And we keep doing this recursively – over time this helps us to learn sequences!
- ◆ btw, this is still not completely describing Q-learning, you'll cover more in 4360

52

## Example

- ◆ Mataric's *Reinforcement Learning in the Multi-Robot Domain* (paper in Stuff of the Week – paper is not examinable, but you should understand this example at the level at which these notes present it!)
- ◆ This paper nicely discusses some of the issues we've already brought up
- ◆ Uses several behaviour-based robots in a puck-gathering setting (think of waste cleanup)
- ◆ Individuals able to deal with the following behaviours:

53

## Behaviours

- ◆ Safe-wandering: moving about without colliding with objects (including other robots)
- ◆ Dispersion: achieve and maintain distance between all robots within sensing range of one another
- ◆ Resting: keeps robot parked for fixed period of time (intended for recharging)
- ◆ Homing: go to particular location

54

## General idea

- ◆ Learning the appropriate conditions for triggering these behaviours
- ◆ Can perceive conditions:
  - Have puck?
  - At home?
  - Near intruder? (distance between robot and nearest neighbour – "personal space")
  - Night-time? (internal clock, desire to recharge)
  - These behaviours and predicates constitute the learning space. *Other behaviours are present but not learned*

55

## Learning Approach

- ◆ It had been thought in the RL community that minimizing reinforcement as much as possible was good, because it diminishes experimenter bias
- ◆ In this work, Mataric was trying to illustrate that it's more important to get power out of learning by shaping reinforcement to suit the task
- ◆ Makes sense for human learning anyway!
- ◆ So she compared learning with a simple reinforcement function to some more complex ones that were tailored more closely to the situation

56

## Learning Approach

- ◆ similar in spirit to Q-learning
- ◆ we have a matrix representing a weighted mapping between percepts and behaviours
- ◆  $A(c,b)$  entry in the matrix is the normalized sum of reinforcement  $R$  received for condition/behaviour pair over time  $T$  (i.e. *average goodness over time*)
- ◆ possible to give positive reinforcement for these events: grasping a puck, dropping puck at home, waking up at home
- ◆ Negative reinforcement for dropping puck away from home, waking up away from home

57

## Events

- ◆ Selecting behaviours is induced by events
- ◆ These may be external: getting in the way of one another, dropping a puck
- ◆ These may be internal (clock signalling night-time, or a wake-up)
- ◆ May be triggered by two *progress indicator* functions
  - one to see if we are moving away from an intruder in our personal space
  - Another to tell if we are moving away from home or not
  - These can also be used give reinforcement under certain conditions (bad if we have a puck and are moving away from our goal, or if we have an intruder and aren't moving away from them)

58

## Handling Events

- ◆ When an event is detected:
- ◆ Call the reinforcement function to deal with current behaviour-condition pair (deals out reinforcement based on current conditions)
- ◆ We terminate current behaviour
- ◆ We choose another behaviour:
  - Untried one if one is available, Otherwise the best
  - This first encourages exploration (ok in this domain, not so great in others!)
  - This is unusual in that most domains would divide it into best  $x\%$  of the time, random  $y\%$
  - She argues that perception is so noisy that this alone gives us enough randomness

59

## Learning

- ◆ Incremental and continuous over time
- ◆ Table is 64 entries:  $2^4$  conditions \* 4 behaviours
- ◆ Compared reinforcement function to others:
  - Basic single goal reward function (puck home) using Q-learning
  - Heterogeneous reward function using multiple goals, summing reinforcement over time (the collection of all reinforceable elements described earlier, except for the progress estimators)
  - Heterogeneous reward function using this as well as the two progress estimators

60

## Success in Learning

- ◆ Many ways to look at success in learning
- ◆ We could look at overall success at a task (e.g. if we're foraging for pucks, the more pucks we get in the better)
  - indirect measurement of learning
- ◆ There are more direct measures as well
- ◆ For example, if we know what the correct set of actions or decisions (policy) should be, we can look at how much of that policy is learned
  - learning should converge to the correct policy
  - if learning doesn't ever converge, something is not learnable given our learning method
- ◆ Also, how fast do we converge?

61

## In This Example

- ◆ 60 trials, 20 of each strategy, 4 robots
- ◆ Hard to compare – we can't just talk about convergence because the amount of time required depends on external events
- ◆ looked at % of correct policy (constructed by hand from observation) robots learned in 15 minutes
- ◆ Q: ~28%
- ◆ R over time: ~52%
- ◆ R over time + PE = 89%

62