

Relational Algebra

COMP 3380 - Databases: Concepts and Usage

Department of Computer Science
The University of Manitoba
Fall 2006

Relational Query Languages

- ❖ **Query languages (QLs):** Allow manipulation and retrieval of data from a DB
- ❖ Relational model supports simple & powerful QLs:
 - Strong formal foundation based on logic
 - Allows for much optimization
- ❖ Query languages \neq programming languages
 - QLs not intended to be used for complex calculations
 - QLs support easy & efficient access to large data sets

Formal Relational Query Languages

- ❖ Two mathematical query languages form the basis for “real” languages (e.g., SQL), and for implementation:
 - **Relational Algebra:** User specifies what data is required & *how to get those data*
 - Procedural (or operational)
 - Very useful for representing execution plans
 - **Relational Calculus:** User specifies what data is required *without specifying how to get those data*
 - Declarative (i.e., non-procedural)

Preliminaries

- ❖ A query is applied to *relation instances*, and **the result of a query is also a *relation instance***
 - Schemas of input relations for a query are fixed
 - The schema for the result of a given query is fixed
 - Determined by definition of query language constructs
- ❖ Positional *vs.* named-attribute notation:
 - **Positional notation** easier for formal definitions, **named-attribute notation** more readable.
 - Both used in SQL

Relational Algebra

❖ Basic operations:

- **Selection (σ):** Selects a subset of rows from a relation
- **Projection (π):** Extracts wanted columns from a relation
- **Cartesian-product (\times):** Allows us to combine two relations
- **Set-difference ($-$):** Returns the tuples that are in R1 but not in R2
- **Union (\cup):** Returns the tuples that are in R1 or R2

Relational Algebra

❖ Additional operations:

- **Intersection (\cap)**
- **Join (\bowtie)**
- **Division (\div)**
- **Rename (ρ)**
- **Assignment (\leftarrow)**

Not essential, but (very) useful

- ❖ Since each operation returns a relation, operations can be *composed* (Algebra is “closed”)

Selection

- ❖ **Selects rows** that satisfy the *selection condition*
- ❖ $\sigma_p(R) = \{t \mid t \in R \text{ and } p(t)\}$ where p is a formula in propositional calculus consisting of terms connected by \wedge , \vee , or \neg . Each term is either
 - $\text{attr}_1 \text{ op attr}_2$, or
 - attr op constant ,where op is one of $=, \neq, >, \geq, <, \leq$

Selection

- ❖ **No duplicates in result**
- ❖ *Schema* of result identical to schema of the (only) input relation
- ❖ *Result* relation can be the input for another relational algebra operation (*Operator composition*)
- ❖ E.g., consider *Acct* (*acctNum*, *branchName*, *bal*). Find all those tuples of *Acct* about the Vancouver branch:

$$\sigma_{\text{branchName}='Vancouver'}(\text{Acct})$$

Example: Selection

❖ Relation R :

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

❖ $\sigma_{A=B \wedge D > 5}(R)$:

A	B	C	D
α	α	1	7
β	β	23	10

Projection

- ❖ **Projects columns**, i.e., extracts only attributes that are in the *projection list*
- ❖ $\pi_{A_1, A_2, \dots, A_k}(R)$ where A_1, A_2, \dots, A_k are attributes in the relation R
- ❖ Useful when we do not want to retrieve unnecessary data, thereby reducing the query cost and/or hiding irrelevant data from the user
- ❖ Projection operator has to **eliminate duplicates**, since relations are *sets*

Projection

- ❖ *Schema* of result contains exactly the attributes in the projection list, with the same names that they had in the (only) input relation.
- ❖ E.g., consider *Acct* (*acctNum*, *branchName*, *bal*). List the account number and balance for each account (i.e., do not list the *branchName* attribute of *Acct*):

$$\pi_{acctNum, bal} (Acct)$$

Example: Projection

- ❖ Relation *R*:

A	B	C
α	10	1
α	20	1
β	30	1
β	40	2

- ❖ $\pi_{A,C}(R)$:

A	C
α	1
α	1
β	1
β	2

 \rightarrow

A	C
α	1
β	1
β	2

Union

- ❖ $R \cup S$ takes two input relations R and S , which must be **union-compatible**
 - (i.e., same number of attributes &
 - “corresponding” attributes have the same domain), &returns a relation instance containing all tuples that occur in R or S
- ❖ $R \cup S = \{t \mid t \in R \text{ or } t \in S\}$
- ❖ Schema of result identical to schema of input relations
- ❖ E.g., consider *Depositor* (*custName*, *acctNum*) & *Borrower* (*custName*, *acctNum*). Find the names of all customers with an account or a loan (i.e., either an account or a loan, or both):
$$\pi_{\text{custName}}(\text{Depositor}) \cup \pi_{\text{custName}}(\text{Borrower})$$

Example: Union

- ❖ Relations R, S :

A	B
α	1
α	2
β	1

R

A	B
α	2
β	3

S

- ❖ $R \cup S$:

A	B
α	1
α	2
β	1
β	3

Intersection

- ❖ $R \cap S$ takes two input relations R and S , which must be **union-compatible**, & returns a relation instance containing all tuples that occur in both R and S
- ❖ $R \cap S = \{t \mid t \in R \text{ and } t \in S\}$
- ❖ Schema of result identical to schema of input relations
- ❖ E.g., consider *Depositor* (*custName*, *acctNum*) & *Borrower* (*custName*, *acctNum*). Find the names of all customers with both an account and a loan.

COMP 3380 (Fall 2006), Leung

15

Example: Intersection

- ❖ Relations R , S :

A	B
α	1
α	2
β	1

R

A	B
α	2
β	3

S

- ❖ $R \cap S$:

A	B
α	2

COMP 3380 (Fall 2006), Leung

16

Set Difference

- ❖ $R - S$ takes two input relations R and S , which must be **union-compatible**, & returns a relation instance containing all tuples that occur in R but not in S
- ❖ $R - S = \{t \mid t \in R \text{ and } t \notin S\}$
- ❖ Schema of result identical to schema of input relations
- ❖ E.g., consider *Depositor* (*custName*, *acctNum*) & *Borrower* (*custName*, *acctNum*). Find the names of all customers with an account but not a loan.

COMP 3380 (Fall 2006), Leung

17

Example: Set Difference

- ❖ Relations R, S :

A	B
α	1
α	2
β	1

R

A	B
α	2
β	3

S

- ❖ $R - S$:

A	B
α	1
β	1

COMP 3380 (Fall 2006), Leung

18

Cartesian Product (or Cross Product)

- ❖ Pairs each row of R with each row of S
- ❖ $R \times S = \{t \ q \mid t \in R \text{ and } q \in S\}$
- ❖ *Result schema* has 1 attr per attr of R and S , with attribute names “inherited” if possible.
- ❖ If $\text{schema}(R)$ and $\text{schema}(S)$ are not disjoint (\rightarrow ambiguous attribute names):
 - **prefix attributes with the relation names, or**
 - **rename the attributes.**

Example: Cartesian Product

- ❖ Relations R, S :

A	B
α	1
β	2

R

A	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

S

- ❖ $T = R \times S$:

$R.A$	B	$S.A$	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

Rename

- ❖ Allows us to name (& refer to) the results of RA expressions; allows us to refer to a relation by more than one name.
- ❖ Avoids ambiguity when the same relation appears in a query more than once
- ❖ $\rho_Y(E)$ returns the expression E under the name Y
- ❖ E.g., $R \bowtie \rho_{R'}(R)$
- ❖ E.g., $\rho_{T(A, B, C, D, E)}(R \bowtie S)$
- ❖ NOTE: Formal RA does not have names for relations

Join (or Natural Join)

- ❖ $R \bowtie S$ joins tuples in relations R and S
- ❖ An equi-join in which equalities are specified on *all* attrs having the same name in R & S
- ❖ *Result schema* similar to Cartesian-product, but only *one* copy of attributes for which equality is specified
- ❖ Fewer tuples than Cartesian-product, might be able to compute more efficiently

Join (or Natural Join)

- ❖ E.g., Consider $R(A, B, C, D)$ & $S(B, D, E)$.
 - Schema of $T = R \bowtie S$ is (A, B, C, D, E)
 - $T = R \bowtie S$

$$= \pi_{A, R.B, C, R.D, E} (\sigma_{R.B=S.B \wedge R.D=S.D} (R \times S))$$
- ❖ NOTE: If there are *no common attributes* in relations P and Q , then $P \bowtie Q = P \times Q$.

Example: Join (1)

- ❖ Relations R, S :

A	B	C	D
α	1	α	a
α	1	γ	a
δ	2	β	b

R

B	D	E
1	a	α
1	a	γ
2	b	δ

S

- ❖ $T = R \bowtie S$:

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ

Example: Join (2)

❖ Relations R, S :

A	B	C	D
α	1	α	a
β	2	γ	a
γ	4	β	b
α	1	γ	a
δ	2	β	b

R

B	D	E
1	a	α
3	a	β
1	a	γ
2	b	δ
3	b	ϵ

S

❖ $T = R \bowtie S$:

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ

Relaxations of (Natural) Join: Equi-Join & Condition Join

❖ **Join (or natural join):**

- An equi-join in which equalities are specified on *all* attributes having the same name in R & S

❖ **Equi-join:**

- A special case of condition join where *cond* contains only **equalities**
- E.g., $U = R \bowtie_B S = \pi_{A, R.B, C, R.D, S.D, E} (\sigma_{R.B=S.B} (R \times S))$

❖ **Condition join (aka theta-join):**

- $R \bowtie_{\text{cond}} S = \sigma_{\text{cond}} (R \times S)$
- E.g., $V = R \bowtie_{R.B < S.B} S = \pi_{A, R.B, C, R.D, S.B, S.D, E} (\sigma_{R.B < S.B} (R \times S))$

Extensions of (Inner) Join: Outer Joins

- ❖ Extension of (inner) joins that avoids loss of information.
- ❖ Compute the **inner join** & then add tuples from one relation that do not match tuples in the other relation to the result of the join.
- ❖ Variations:
 - **Left outer join** ($\bowtie\leftarrow$)
 - **Right outer join** ($\rightarrow\bowtie$)
 - **(Full) outer join** ($\bowtie\leftarrow\rightarrow$)

COMP 3380 (Fall 2006), Leung 27

27

Example: Outer Joins

❖ Relations R, S :

A	B
α	1
β	2
γ	3

R

A	C
β	b
γ	c
δ	d

S

❖ $R \bowtie S$:

A	B	C
β	2	b
γ	3	c

$R \bowtie S$:

A	B	C
α	1	NULL
β	2	b
γ	3	c

$R \ltimes S$:

A	B	C
β	2	b
γ	3	c
δ	NULL	d

$R \ltimes S$:

A	B	C
α	1	NULL
β	2	b
γ	3	c
δ	NULL	d

COMP 3380 (Fall 2006), Leung

28

A	C
β	b
γ	c
δ	d

S

A	B	C
α	1	NULL
β	2	b
γ	3	c
δ	NULL	d

28

Division

- ❖ Suited to queries that include the phrase “for all” (e.g., “find boats that are reserved by **all** sailors”)
- ❖ $R \div S = \{t \mid t \in \pi_{\text{schema}(R-S)}(R) \wedge \forall u \in S (tu \in R)\}$
 - $R \div S = \{ \langle x \rangle \mid \forall y \in S (\exists \langle x, y \rangle \in R) \}$
 - i.e., $R \div S$ contains all x tuples (boats in the answer) such that for *every* y tuple (every sailor) in S , there is an xy tuple (a reservation made by sailor y on boat x) in R
 - i.e., If the set of y values (sailors) associated with an x value (boat) in R contains all y values in S (i.e., boat x reserved by all sailors), then the x value is in $R \div S$
- ❖ In general, x and y can be any lists of attrs; y is the list of attrs in S , and $x \cup y$ is the list of attrs of R .

Division

- ❖ Division is not essential op; just a useful shorthand.
 - (Also true of joins, but joins are so common that systems implement joins specially)
- ❖ Idea: For $R \div S$, compute all x values that are **not “disqualified”** by some y value in S .
 - x value is *disqualified* if by attaching y value from S , we obtain an xy tuple that is not in R .
 - $R \div S = \pi_x(R) - \pi_x((\pi_x(R) \times S) - R)$

Example: Division (1)

- ❖ E.g., consider *Sailors*(sID, sName, rating, age) & *Reserves*(sID, bID, rDate).

Find IDs of the boats who have reserved all sailors:

- $$\pi_{sID, bID}(Reserves) \div \pi_{sID}(Sailors)$$
- $$= \{t \mid t \in \pi_{bID}(Reserves) \wedge \forall u \in \pi_{sID}(Sailors) (tu \in \pi_{sID, bID}(Reserves))\}$$
- $$= \{ \langle bID \rangle \mid \forall sID \in \pi_{sID}(Sailors) (\exists \langle sID, bID \rangle \in \pi_{sID, bID}(Reserves)) \}$$
- ❖ The result contains the ID of each boat x such that for *every* sailor y in *Sailors*, there is a reservation on boat x made by sailor y in *Reserves*.
 - ❖ If the set of sIDs associated with any bID in *Reserves* contains **all** the sIDs in *Sailors* (i.e., a boat reserved by all sailors), then such a bID is in the result of the division $\pi_{sID, bID}(Reserves) \div \pi_{sID}(Sailors)$.
 - ❖ In other words, bID is *disqualified* if by attaching sID from *Sailors*, we obtain a $\langle sID, bID \rangle$ tuple that is not in *Reserves*:
- $$\pi_{sID, bID}(Reserves) \div \pi_{sID}(Sailors)$$
- $$= \pi_{bID}(Reserves) - \pi_{bID}((\pi_{sID}(Sailors) \times \pi_{bID}(Reserves)) - \pi_{sID, bID}(Reserves))$$

Example: Division (2)

- ❖ E.g., consider *Boats*(bID, bName, color) & *Reserves*(sID, bID, rDate).

Find IDs of the sailors who have reserved all boats:

- $$\pi_{sID, bID}(Reserves) \div \pi_{bID}(Boats)$$
- $$= \{t \mid t \in \pi_{sID}(Reserves) \wedge \forall u \in \pi_{bID}(Boats) (tu \in \pi_{sID, bID}(Reserves))\}$$
- $$= \{ \langle sID \rangle \mid \forall bID \in \pi_{bID}(Boats) (\exists \langle sID, bID \rangle \in \pi_{sID, bID}(Reserves)) \}$$
- ❖ The result contains the ID of each sailor x such that for *every* boat y in *Boats*, there is a reservation made by sailor x on boat y in *Reserves*.
 - ❖ If the set of bIDs associated with any sID in *Reserves* contains **all** the bIDs in *Boats* (i.e., a sailor reserved all boats), then such an sID is in the result of the division $\pi_{sID, bID}(Reserves) \div \pi_{bID}(Boats)$.
 - ❖ In other words, sID is *disqualified* if by attaching bID from *Boats*, we obtain a $\langle sID, bID \rangle$ tuple that is not in *Reserves*:
- $$\pi_{sID, bID}(Reserves) \div \pi_{bID}(Boats)$$
- $$= \pi_{sID}(Reserves) - \pi_{sID}((\pi_{sID}(Reserves) \times \pi_{bID}(Boats)) - \pi_{sID, bID}(Reserves))$$

Example: Division (3)

sno	pno	pno	pno	pno
s1	p1	p2	p2	p1
s1	p2		p4	p2
s1	p3			p4
s1	p4			
s2	p1			
s2	p2			
s3	p2			
s4	p2			
s4	p4			

R

pno
p2

$S1$

sno
s1
s2
s3
s4

$R \div S1$

pno
p2
p4

$S2$

sno
s1
s4

$R \div S2$

pno
p1
p2
p4

$S3$

sno
s1

$R \div S3$

COMP 3380 (Fall 2006), Leung

33

Example: Division (4)

❖ Relations R, S :

A	B
α	1
α	2
α	3
β	1
β	2
γ	1
δ	1
δ	3
δ	4
ϵ	6
ϵ	1

R

B
1
2

S

❖ $R \div S$:

A
α
β

Group by **A**,
find the groups having all **B**

COMP 3380 (Fall 2006), Leung

34

Example: Division (5)

❖ Relations R, S :

A	B	C	D	E
α	a	α	a	1
α	a	γ	a	1
α	a	γ	b	1
β	a	γ	a	1
β	a	γ	b	3
γ	a	γ	a	1
γ	a	γ	b	1
γ	a	β	b	1

R

D	E
a	1
b	1

S

Group by **ABC**,
find the grps having all **DE**

❖ $R \div S$:

A	B	C
α	a	γ
γ	a	γ

Examples: Division (3-5)

- ❖ E.g., Consider $R(sno, pno)$ & $S(pno)$.
 - Schema of $T = R \div S$ is (**sno**)
 - $T = R \div S = \pi_{sno}(R) - \pi_{sno}((\pi_{sno}(R) \times S) - R)$
- ❖ E.g., Consider $R(A, B)$ & $S(B)$.
 - Schema of $T = R \div S$ is (**A**)
 - $T = R \div S = \pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$
- ❖ E.g., Consider $R(A, B, C, D, E)$ & $S(D, E)$.
 - Schema of $T = R \div S$ is (**A, B, C**)
 - $T = R \div S = \pi_{A,B,C}(R) - \pi_{A,B,C}((\pi_{A,B,C}(R) \times S) - R)$

More Examples

- ❖ Consider the following 3 relations:
 - *Sailors* (sID, sName, rating, age)
 - *Boats* (bID, bName, color)
 - *Reserves* (sID, bID, rDate)
- ❖ *S1* is an instance of *Sailor*, *R1* is an instance of *Reserves*
- ❖ Retrieve rows about sailors with rating > 8:
$$\sigma_{rating > 8}(S1)$$
- ❖ Find all sailor names & ratings:
$$\pi_{sName, rating}(S1)$$
- ❖ Find only the ages of sailors:
$$\pi_{age}(S1)$$

More Examples

- ❖ Compute the names & ratings of 8⁺-rated sailors:
$$\pi_{sName, rating}(\sigma_{rating \geq 8}(S1))$$
- ❖ Suppose there is another instance of *Sailor* (say, *S2*)
- ❖ Retrieve rows about sailors in either *S1* or *S2*:
$$S1 \cup S2$$
- ❖ Retrieve rows about sailors in both *S1* and *S2*:
$$S1 \cap S2$$
- ❖ Retrieve rows about sailors in *S1* but not in *S2*:
$$S1 - S2$$

More Examples

- ❖ Pair each tuple of sailors with each tuple of reserves:

$$S1 \times R1$$

- ❖ Find all sailors who have reserved boats:

$$S1 \bowtie R1$$

- ❖ Find names of sailors who have reserved boat #103:

$$\pi_{sName} (\sigma_{bID=103} (S1 \bowtie R1))$$

- ❖ Find names of sailors who have reserved a red boat:

$$\pi_{sName} (\sigma_{color='red'} (B1) \bowtie R1 \bowtie S1)$$

More Examples

- ❖ Find names of sailors who have reserved a red or a green boat:

$$\pi_{sName} (\sigma_{color='red' \vee color='green'} (B1) \bowtie R1 \bowtie S1)$$

or

$$\pi_{sName} (\pi_{sID} ([\sigma_{color='red'} (B1) \cup \sigma_{color='green'} (B1)] \bowtie R1) \bowtie S1)$$

or

$$\pi_{sName} (S1 \bowtie [\pi_{sID} (\sigma_{color='red'} (B1) \bowtie R1) \cup \pi_{sID} (\sigma_{color='green'} (B1) \bowtie R1)])$$

More Examples

- ❖ Find names of sailors who have reserved a red and a green boat:

$$\pi_{sName} (S1 \bowtie [\pi_{sID} (\sigma_{color='red'} (B1) \bowtie R1) \cap \pi_{sID} (\sigma_{color='green'} (B1) \bowtie R1)])$$

How about the following?

$$\pi_{sName} (\sigma_{color='red' \wedge color='green'} (B1) \bowtie R1 \bowtie S1)$$

How about the following?

$$\pi_{sName} (\pi_{sID} [\sigma_{color='red'} (B1) \cap \sigma_{color='green'} (B1)] \bowtie R1 \bowtie S1)$$

More Examples

- ❖ Find names of sailors who have reserved a red but not a green boat:

$$\pi_{sName} (S1 \bowtie [\pi_{sID} (\sigma_{color='red'} (B1) \bowtie R1) - \pi_{sID} (\sigma_{color='green'} (B1) \bowtie R1)])$$

How about the following?

$$\pi_{sName} (\sigma_{color='red' \wedge color \neq 'green'} (B1) \bowtie R1 \bowtie S1)$$

How about the following?

$$\pi_{sName} (\pi_{sID} [\sigma_{color='red'} (B1) - \sigma_{color='green'} (B1)] \bowtie R1 \bowtie S1)$$

More Examples

- ❖ Find IDs of sailors who have reserved all boats:

$$\pi_{sID, bID} (R1) \div \pi_{bID} (B1)$$

- ❖ Find names of sailors who have reserved all boats:

$$\pi_{sName} (S1 \bowtie (\pi_{sID, bID} (R1) \div \pi_{bID} (B1)))$$

Modification of the DB

- ❖ The content of the DB may be modified using the following operations:
 - Insertion
 - Deletion
 - Update
- ❖ All these operations can be expressed using the **assignment operator** (\leftarrow)

Insertion

- ❖ To insert data into a relation, we either:
 - specify a tuple to be inserted, or
 - write a query whose result is a set of tuples to be inserted
- ❖ In RA, an insertion is expressed by:
$$R \leftarrow R \cup E$$
where R is a relation and E is a RA expression.
- ❖ The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple.

Example: Insertion

- ❖ Consider $Acct(acctNum, branchName, bal)$ & $Depositor(custName, acctNum)$.
- ❖ Insert information into the DB specifying that Smith has \$1200 in account A973 at the Vancouver branch.

$$Acct \leftarrow Acct \cup \{('A973', 'Vancouver', 1200)\}$$
$$Depositor \leftarrow Depositor \cup \{('Smith', 'A973')\}$$

Deletion

- ❖ A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database.
- ❖ Can delete only *whole tuples*; cannot delete values on only particular attributes
- ❖ A deletion is expressed in RA by:

$$R \leftarrow R - E$$

where R is a relation and E is a RA query.

Example: Deletion

- ❖ Consider $Acct (acctNum, branchName, bal)$ & $Loan (loanNum, branchName, amt)$.
- ❖ Delete all account records in the Toronto branch

$$Acct \leftarrow Acct - \sigma_{branchName = 'Toronto'}(Acct)$$

- ❖ Delete all loan records with amount in the range of \$0 to \$50

$$Loan \leftarrow Loan - \sigma_{amt \geq 0 \wedge amt \leq 50}(Loan)$$

Update

- ❖ A mechanism to change a value in a tuple without changing *all* values in the tuple
- ❖ $R \leftarrow \pi_{F_1, F_2, \dots, F_k}(R)$
- ❖ Each F_i is either
 - the i -th attribute of R (if the i -th attribute is not updated), or
 - an expression, involving only constants and the attributes of R , which gives the new value for the attribute F_i (if the attribute F_i is to be updated)

Example: Update

- ❖ Consider $Acct(acctNum, branchName, bal)$.
- ❖ Make interest payments by increasing all balances by 5%

$$Acct \leftarrow \pi_{acctNum, branchName, bal*1.05}(Acct)$$

- ❖ Other Notation: $\delta_{bal \leftarrow bal*1.05}(Acct)$

Example: Update

- ❖ Consider $Acct(acctNum, branchName, bal)$.
- ❖ Pay all accounts with balances over \$10,000 a 6% interest and pay all others 5%

$$Acct \leftarrow \pi_{acctNum, branchName, bal*1.06} (\sigma_{bal>10000} (Acct)) \\ \cup \pi_{acctNum, branchName, bal*1.05} (\sigma_{bal\leq 10000} (Acct))$$

- ❖ Other Notation:

$$\delta_{bal \leftarrow bal*1.06} (\sigma_{bal>10000} (Acct)) \\ \delta_{bal \leftarrow bal*1.05} (\sigma_{bal\leq 10000} (Acct))$$

Summary

- ❖ The relational model has rigorously defined query languages that are simple and powerful.
- ❖ Relational algebra (RA) is more operational; useful as internal representation for query evaluation plans.
- ❖ *Several ways* of expressing a given query; a query optimizer should choose the most efficient version.