

The Relational Model

COMP 3380 - Databases: Concepts and Usage

Department of Computer Science
The University of Manitoba
Fall 2006

Review: Conceptual/Logical Data Model

❖ Examples:

1. Entity-Relationship model

- **Entities** are relations (i.e., tables with rows & column)
- **Relationships** describe association between entities

2. Relational model - *most widely used model*

- **Relations** are basically tables with rows & columns

Why Study the Relational Model?

- ❖ Proposed by E. F. Codd (1923-2003) in 1970
- ❖ Most widely used model
 - Vendors: IBM (DB2), Informix, Microsoft (Access), Oracle, Sybase, etc.
- ❖ “Legacy systems” in older models
 - E.g., IBM’s IMS
- ❖ Recent competitor: Object-oriented model
 - E.g., ObjectStore, Versant, Ontos
 - A synthesis emerging: **Object-relational model** (e.g., Informix Universal Server, UniSQL, O2, Oracle, DB2)

Relational DB

- ❖ **Relational database:** A set of *relations*
- ❖ **Relation:** Made up of 2 parts
 1. **Instance:** A *table*, with rows & columns
 - #rows = cardinality, #columns = degree or arity
 2. **Schema:** Basic info describing a table or relation (i.e., specifies name of relation, plus name & type of each column)
 - E.g., *Students* (*sID*: string, *sName*: string, *loginID*: string)
- ❖ Can think of a relation instance as a ***set of rows*** (i.e., **all rows/ tuples/ records are distinct**)

Example: Instance of Students Relation

Students

sID	sName	loginID
3666	Adam	adam@cs
3688	Ben	ben@ece

- ❖ Cardinality (i.e., #rows or #tuples) = 2
- ❖ Degree (aka arity; i.e., #cols or #attrs or #fields) = 3
- ❖ All rows are distinct
- ❖ **Order of rows is unimportant**

Relational DB: Relation Instance

- ❖ Given n domain sets D_1, D_2, \dots, D_n , a **relation instance** r is
 - a subset of $D_1 \times D_2 \times \dots \times D_n$
 - a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$
- ❖ A domain D_i
 - is the set of allowed values for each attribute a_i (normally required to be *atomic*, i.e., not multi-valued & not composite)
 - contains a special value **NULL** indicating that the value is not known (or is non-applicable)

Example: Relation Instance

- ❖ If $custName = \{\text{Adam, Ben, Charles, Don}\}$,
 $custStreet = \{\text{Robson, St-Catherine, Yonge}\}$, &
 $custCity = \{\text{Montréal, Toronto, Vancouver}\}$,
then, *Customer* $r = \{ (\text{Adam, St-Catherine, Montréal}),$
 $(\text{Ben, Yonge, Toronto}),$
 $(\text{Charles, Yonge, Toronto}),$
 $(\text{Don, Robson, Vancouver}) \}$ is a relation instance
over $custName \times custStreet \times custCity$

Relational DB: Relation Schema

- ❖ If A_1, \dots, A_n are attributes with domains D_1, \dots, D_n , then $R(A_1:D_1, \dots, A_n:D_n)$ is a **relation schema** that defines a relation type
- ❖ E.g., *Customer* ($custName: string,$
 $custStreet: string,$
 $custCity: string$)
- ❖ **Normalization theory** deals with how to design relation schemas
- ❖ A DB consists of multiple relations

Creating Relations

- ❖ Data definition language (DDL), such as SQL, for creating Students relation

```
CREATE TABLE Students  
  (sID CHAR(20),  
   sName CHAR(20),  
   loginID CHAR(10))
```

- ❖ Creates Enrolled relation

```
CREATE TABLE Enrolled  
  (sID CHAR(20),  
   cID CHAR(20),  
   grade CHAR(2))
```

Destroying and Altering Relations

- ❖ Destroys the relation Students. The schema information *and* the tuples are deleted.

```
DROP TABLE Students
```

- ❖ Alters the schema of Students by adding a new column (called firstYear); every tuple in the current instance is extended with a NULL value in the new column.

```
ALTER TABLE Students  
  ADD firstYear INTEGER
```

Adding Tuples

- ❖ Data manipulation language (DML), such as SQL, for inserting a single tuple

```
INSERT INTO Students (sID, sName, loginID)
VALUES ('3666', 'Adam', 'adam@cs')
```

or

```
INSERT INTO Students
VALUES ('3666', 'Adam', 'adam@cs')
```

Adding Tuples

- ❖ Inserts tuples based on the result of a query

```
INSERT INTO Freshmen (sID, sName, loginID)
SELECT sID, sName, loginID
FROM Students
WHERE firstYear = 1
```

Deleting and Modifying Tuples

- ❖ Deletes all tuples satisfying some conditions (e.g., name = 'Adam')

```
DELETE FROM Students S
WHERE S.sName = 'Adam'
```

- ❖ Changes attributes (e.g., grade) of all tuples satisfying some conditions (e.g., sID='3666' & cID='comp3380')

```
UPDATE Enrolled E
SET E.grade = 'A'
WHERE E.sID = '3666' and E.cID = 'comp3380'
```

Integrity Constraints (ICs)

- ❖ **Integrity constraints (ICs)** are conditions that must be true for *any* instance of the database (e.g., domain constraints)
 - **ICs are specified when schema is defined**
 - **ICs are checked when relations are modified** (e.g., insertions/ deletions/ updates that violate ICs are disallowed)
- ❖ A **legal instance** of a relation is one that satisfies all specified ICs
 - DBMS should not allow illegal instances

Integrity Constraints (ICs)

- ❖ ICs can be used to ensure application semantics (e.g., *sID* is a key), or to prevent inconsistencies (e.g., *sName* has to be a string, *age* must be < 150)
- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning
 - Avoids data entry errors, too!


Types of Integrity Constraints (ICs)

- ❖ **Domain constraints**
 - Attribute values must be of right type. Always enforced.
- ❖ **Key constraints** (aka **primary key constraints**)
 - Forbid duplicate primary key values (i.e., no 2 tuples have the same primary key value; *primary key value is unique*)
- ❖ **Entity integrity constraints**
 - Primary key is NOT NULL
- ❖ **Referential integrity constraints** (aka **foreign key constraints**)
 - Ensure consistency of tuple references across tables (i.e., no dangling references)
- ❖ Other general constraints (e.g., semantics)

Possible Violations of Integrity Constraints (ICs)

- ❖ Recall: ICs are specified when schema is defined; ICs are checked when relations are modified (e.g., insertions, deletions, updates)
- ❖ Insertions
 - may violate (1) domain constraints, (2) key constraints (primary key constraints), (3) entity integrity constraints, & (4) referential integrity constraints (foreign key constraints)
- ❖ Deletions
 - may violate (4) referential integrity constraints
- ❖ Updates (~ deletions follow by insertions)
 - may violate (1) domain constraints, (2) key constraints, (3) entity integrity constraints, & (4) referential integrity constraints

Review: Primary Key Constraints



- ❖ **Superkey**
 - A set of one or more attributes that (taken collectively) identify uniquely an entity in an entity set
 - E.g., {sID}, {sID, sName}, {sID, sName, loginID}, {sID, loginID}, {loginID}, {loginID, sName}
- ❖ **Candidate key**
 - Minimal superkey (i.e., a superkey for which no proper subset is a superkey)
 - E.g., {sID}, {loginID} **More than 1 candidate keys!**
- ❖ **Primary key**
 - The candidate key chosen as the principal means to identify entities in an entity set
 - E.g., {sID}

Primary and Candidate Keys

- ❖ Possibly many **candidate keys** (specified using **UNIQUE**), one of which is chosen as the **primary key**.

```
CREATE TABLE Students
(sID CHAR(20) NOT NULL,      ← For some DBMS, this
sName CHAR(20),             "NOT NULL" is optional
loginID CHAR(10) NOT NULL,
PRIMARY KEY (sID),
UNIQUE (loginID))
```

Foreign Keys & Referential Integrity

- ❖ **Foreign key**: Set of attributes in one relation (i.e., *referencing relation*) that is used to “refer” to a tuple in another relation (i.e., *referenced relation*).
 - Must correspond to **candidate key of the referenced relation**. Like a “logical pointer”.
- ❖ E.g., sID is a foreign key referring to Students
 - Enrolled (sID: string, cID: string, grade: string)
 - If all foreign key constraints are enforced, **referential integrity** is achieved (i.e., no dangling references).
 - Can you name a data model without referential integrity?
 - Links in HTML!

Foreign Keys

- ❖ Only students listed in the Students relation should be allowed to enrol for courses.

```
CREATE TABLE Enrolled  
  (sID CHAR(20) NOT NULL,  
   cID CHAR(20) NOT NULL,  
   grade CHAR(2),  
   PRIMARY KEY (sID, cID),  
   FOREIGN KEY (sID) REFERENCES Students,  
   FOREIGN KEY (cID) REFERENCES Courses)
```

Enforcing Referential Integrity

- ❖ Consider Students and Enrolled; *sID* in Enrolled is a foreign key that references Students.
- ❖ What should be done if an Enrolled tuple with a non-existent student ID is inserted?
 - Reject it!

Enforcing Referential Integrity

- ❖ What should be done if a Students tuple is deleted?
 - **Default:** Disallow deletion of a Students tuple that is referred to.
 - **Other Options:**
 - Also delete all Enrolled tuples that refer to it (via the use of **ON DELETE CASCADE**).
 - Set *sID* in Enrolled tuples that refer to it to a *default sID* (via the use of **ON DELETE SET DEFAULT**).
 - In SQL: Set *sID* in Enrolled tuples that refer to it to a special value NULL (via the use of **ON DELETE SET NULL**).
- ❖ Similar if primary key of Students tuple is updated.

Examples 1 & 2: Referential Integrity


- ❖ Default (i.e., ON DELETE NO ACTION):

```
CREATE TABLE Enrolled
(sID CHAR(20) NOT NULL,
cID CHAR(20) NOT NULL,
grade CHAR(2),
PRIMARY KEY (sID, cID),
FOREIGN KEY (sID) REFERENCES Students,
FOREIGN KEY (cID) REFERENCES Courses)
```

- ❖ Option 1 - ON DELETE CASCADE:

```
... FOREIGN KEY (sID) REFERENCES Students
ON DELETE CASCADE, ...
```

Example 3: Referential Integrity



❖ Option 2 - ON DELETE SET DEFAULT:

```
CREATE TABLE Manages (
    dID INTEGER NOT NULL,
    eID INTEGER DEFAULT 0,
    since DATE,
    PRIMARY KEY (dID),
    FOREIGN KEY (eID) REFERENCES Employees
        ON DELETE SET DEFAULT,
    FOREIGN KEY (dID) REFERENCES Departments)
```

Example 4: Referential Integrity

❖ Option 3 - ON DELETE SET NULL:

```
CREATE TABLE Manages (
    dID INTEGER NOT NULL,
    eID INTEGER,
    since DATE,
    PRIMARY KEY (dID),
    FOREIGN KEY (eID) REFERENCES Employees
        ON DELETE SET NULL,
    FOREIGN KEY (dID) REFERENCES Departments)
```

Where Do ICs Come From?

- ❖ ICs are based upon the semantics of the real-world enterprise that is being described in the database relations
- ❖ By looking at an instance:
 - We **can check a database instance to see if an IC is violated**
 - We **can NEVER infer that an IC is true**
 - **An IC is a statement about *all possible instances*!**
 - From the example, we know *sName* is not a key, but the assertion that *sID* is a key is given to us.
- ❖ Key and foreign key ICs are the most common; more general ICs supported too.

Views

- ❖ A **view** is just a relation, but we store a *definition* (rather than a set of tuples).

```
CREATE VIEW YoungActiveStudents (sID, sName, grade) AS
  SELECT S.sID, S.sName, E.grade
  FROM Students S, Enrolled E
  WHERE S.sID = E.sID and S.age < 21
```

- ❖ Views can be dropped using the **DROP VIEW** command.

```
DROP VIEW YoungActiveStudents
```

Views and Security


- ❖ Views provides the support for *logical data independence*
- ❖ Views can be used to present necessary information (or a summary), while *hiding details* in underlying relation(s).
 - Given YoungActiveStudents (but not Students or Enrolled), we can find students who are enrolled in some courses without revealing the *cIDs* of the courses they are enrolled in.

Logical DB Design: ER to Relational Model

Review: Database Design Process


1. Requirements collection & analysis
2. Conceptual DB design (ER model)
3. **Logical design (data model mapping, e.g., map ER to tables)**
4. Schema refinement (e.g., normalization)
5. Physical design
6. System implementation & tuning (e.g., application & security design)

Logical DB Design: ER to Relational Model



```
CREATE TABLE Employees
(eID INTEGER NOT NULL,
 eName CHAR(20),
 addr CHAR(100),
 PRIMARY KEY (eID))
```


Logical DB Design: ER to Relational Model




```

CREATE TABLE Manages (
    dID INTEGER NOT NULL,
    eID INTEGER,
    since DATE,
    PRIMARY KEY (dID),
    FOREIGN KEY (eID) REFERENCES Employees,
    FOREIGN KEY (dID) REFERENCES Departments)
    
```

Logical DB Design: ER to Relational Model

- ❖ **One table for each entity set**
 - Attrs are the attrs of the entity set
- ❖ **One table for each relationship set**
 - Attrs include the primary keys of participating entity sets (as foreign keys) & the descriptive attrs of the relationship set
- ❖ Alternative:
 - Include the info about a relationship set (e.g., *Manages*) in the table corresponding to the relevant entity set (e.g., *Dept*)
 - E.g., merge *Dept* (*dID*, *dName*, *budget*) and *Manages* (*dID*, *eID*, *since*) into *DeptMgr* (*dID*, *dName*, *budget*, *eID*, *since*)

Logical DB Design: ER to Relational Model

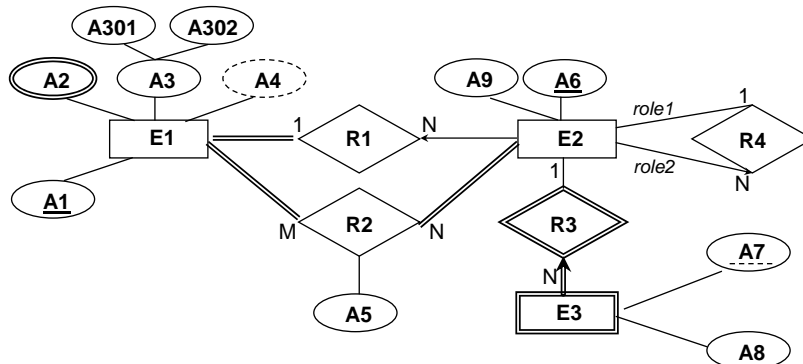


```

CREATE TABLE DependentPolicy (
    eID CHAR(11) NOT NULL, dName CHAR(20) NOT NULL,
    age INTEGER, cost REAL,
    PRIMARY KEY (eID, dName),
    FOREIGN KEY (eID) REFERENCES Employees
    ON DELETE CASCADE)
    
```

Logical DB Design: ER to Relational Model

❖ ER Model:



Logical DB Design: ER to Relational Model

- ❖ Relational Model:
WLOG, let us assume that all attributes are of type CHAR(20)
- ❖ **CREATE TABLE E1 (**
 A1 CHAR(20) **NOT NULL,**
 A301 CHAR(20),
 A302 CHAR(20),
 PRIMARY KEY (A1))
- ❖ **CREATE TABLE E2 (**
 A6 CHAR(20) **NOT NULL,**
 A9 CHAR(20),
 PRIMARY KEY (A6))

Logical DB Design: ER to Relational Model

- ❖ **CREATE TABLE mvA2 (**
 A1 CHAR(20) **NOT NULL,**
 a2 CHAR(20) **NOT NULL,**
 PRIMARY KEY (A1, a2),
 FOREIGN KEY (A1) REFERENCES E1)
- ❖ Note:
 a2 (e.g., author) is a single-valued attribute for storing each of
 the multiple A2 values (e.g., authors)

Logical DB Design: ER to Relational Model

```
❖ CREATE TABLE E3 (  
    A6 CHAR(20) NOT NULL,  
    A7 CHAR(20) NOT NULL,  
    A8 CHAR(20),  
    PRIMARY KEY (A6, A7),  
    FOREIGN KEY (A6) REFERENCES E2  
        ON DELETE CASCADE)
```

Logical DB Design: ER to Relational Model

```
❖ CREATE TABLE R1 (  
    A1 CHAR(20),  
    A6 CHAR(20) NOT NULL,  
    PRIMARY KEY (A6),  
    FOREIGN KEY (A1) REFERENCES E1,  
    FOREIGN KEY (A6) REFERENCES E2)  
❖ CREATE TABLE R2 (  
    A1 CHAR(20) NOT NULL,  
    A6 CHAR(20) NOT NULL,  
    A5 CHAR(20),  
    PRIMARY KEY (A1,A6),  
    FOREIGN KEY (A1) REFERENCES E1,  
    FOREIGN KEY (A6) REFERENCES E2)
```

Logical DB Design: ER to Relational Model

❖ **CREATE TABLE R4 (**
 A6₁ CHAR(20),
 A6₂ CHAR(20) **NOT NULL,**
 PRIMARY KEY (A6₂),
 FOREIGN KEY (A6₁) REFERENCES E2(A6),
 FOREIGN KEY (A6₂) REFERENCES E2(A6))

Logical DB Design: ER to Relational Model

- ❖ Relational Model:
- E1 (A1, A301, A302)
 - E2 (A6, A9)
 - mvA2 (A1, a2) *where* foreign key (A1) references E1
 - E3 (A6, A7, A8) *where* foreign key (A6) references E2
on delete cascade
 - R1 (A1, A6) *where* foreign key (A1) references E1,
foreign key (A6) references E2
 - R2 (A1, A6, A5) *where* foreign key (A1) references E1,
foreign key (A6) references E2
 - R4 (A6₁, A6₂) *where* foreign key (A6₁) references E2(A6),
foreign key (A6₂) references E2(A6)

Logical DB Design: ER to Relational Model

❖ Alternative:

Similar to the *Dept/Manages/DeptMgr* example, E2 can be merged with R1 and R4:

- E2 (A6, A9)
- R1 (A1, A6) where foreign key (A1) references E1,
foreign key (A6) references E2
- R4 (A6₁, A6₂) where foreign key (A6₁) references E2(A6),
foreign key (A6₂) references E2(A6)

to form E2':

- E2' (A6, A9, A1, A6₁) where foreign key (A1) references E1,
foreign key (A6₁) references E2'(A6)

Logical DB Design: ER to Relational Model

❖ Relational Model (after the merge):

- E1 (A1, A301, A302)
- mvA2 (A1, a2) where foreign key (A1) references E1
- E2 (A6, A9, A1, A6₁) where foreign key (A1) references E1,
foreign key (A6₁) references E2(A6)
- E3 (A6, A7, A8) where foreign key (A6) references E2
on delete cascade
- R2 (A1, A6, A5) where foreign key (A1) references E1,
foreign key (A6) references E2

Comments on Logical DB Design: Textbook Approach

1. Mapping of strong entity sets
 2. Mapping of weak entity sets
 3. Mapping of binary 1-to-1 relationship sets
 4. Mapping of binary 1-to-many relationship sets
 5. Mapping of binary many-to-many relationship sets
 6. Mapping of multi-valued attributes
 7. Mapping of n -ary relationship sets
- ➔ *Gives the same relational model!*

Comments on Logical DB Design: Textbook Approach

Example:

1. E1 (A1, A301, A302)
E2 (A6, A9)
2. E3 (A6, A7, A8) where foreign key (A6) references E2
on delete cascade
4. With R1 (A1, A6), E2 becomes:
E2 (A6, A9, A1) where foreign key (A1) references E1
With R4 (A6₁, A6₂), E2 then becomes:
E2 (A6, A9, A1, A6₁) where foreign key (A1) references E1,
foreign key (A6₁) references E2(A6)
5. R2 (A1, A6, A5) where foreign key (A1) references E1,
foreign key (A6) references E2
6. mvA2 (A1, a2) where foreign key (A1) references E1

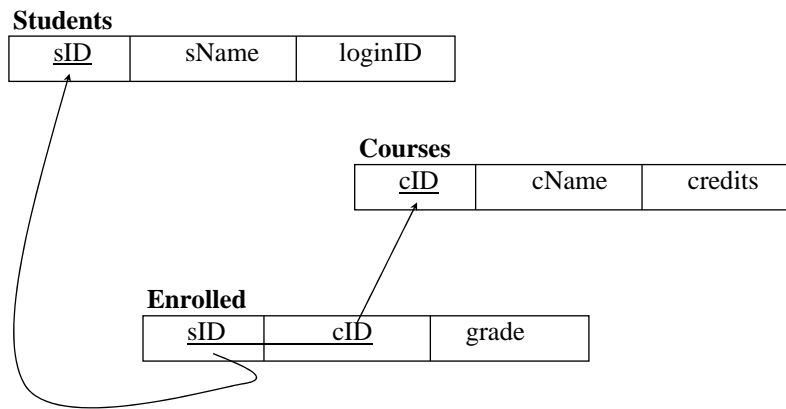
Comments on Logical DB Design: Notation 1

- ❖ **CREATE TABLE** Students
(sID CHAR(20) **NOT NULL**, sName CHAR(20),
loginID CHAR(10), **PRIMARY KEY** (sID))
- ❖ **CREATE TABLE** Courses
(cID CHAR(20) **NOT NULL**, cName CHAR(20),
credits INTEGER, **PRIMARY KEY** (cID))
- ❖ **CREATE TABLE** Enrolled
(sID CHAR(20), cID CHAR(20),
grade CHAR(2)) **PRIMARY KEY** (sID, cID),
FOREIGN KEY (sID) **REFERENCES** Students,
FOREIGN KEY (cID) **REFERENCES** Courses)

Comments on Logical DB Design: Notation 2

- ❖ Students (sID, sName, loginID)
- ❖ Courses (cID, cName, credits)
- ❖ Enrolled (sID, cID, grade) *where*
foreign key (sID) references Students,
foreign key (cID) references Courses

Comments on Logical DB Design: Notation 3



COMP 3380 (Fall 2006), Leung

49

Summary of Relational Model

- ❖ A tabular representation of data
- ❖ Simple & intuitive, currently the most widely used
- ❖ Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
 - Two important ICs: Primary and foreign keys
 - In addition, we *always* have domain constraints
- ❖ Powerful & natural query languages exist

COMP 3380 (Fall 2006), Leung

50

Summary of Relational Model

❖ Rules to translate ER to relational model:

- One table for each entity set
- One table for each multi-valued attribute of the entity set
- One table for each relationship set
 - Except for the identifying relationship set of the weak entity set
- Tables could be merged (if necessary)